

Why the synthesis was needed: each prior model wins only ONE axis
v14 buys content-invariance but loses imitation-catching; v14b keeps both

