

Confusion at deployment threshold — held-out Claude+Gemini eval
green = correct decision, red = error

v12 (multi-gen) — mahalanobis
5% false-alarm operating point

Genuine sender	251 (95%)	13 (5%)
LLM forgery	198 (61%)	129 (39%)
Wrong human	2 (0%)	598 (100%)
	Accepted (allowed)	Flagged (blocked)
	model decision	

v14b (SYNTHESIS) — mahalanobis
5% false-alarm operating point

Genuine sender	251 (95%)	13 (5%)
LLM forgery	40 (12%)	287 (88%)
Wrong human	50 (8%)	550 (92%)
	Accepted (allowed)	Flagged (blocked)
	model decision	